

# Assignment 6

*Veronika Palotai*

*2019 10 22*

**Exercise 0:** Manage your time. If you spend too much time on the assignment, it may be too long, and you may benefit from skipping an exercise. Since this lesson is lost on most people, exercise 0 requires you to skip one of the other exercises this week. That is, write down which exercise you want to skip and why (“saves the most of time because it is the hardest” or “it’s the most useless as you are quite confident in material”). Do not work on it. Of course, if you skip any additional exercise due to lack of time, highlight this, but that will cost some part of the grade, whereas skipping an exercise as part of exercise 0 will not.

**Answer:** I chose to skip **Exercise 4** due to lack of time.

**Exercise 1:** Using the `nycflights13` data. Note that it also contains a tibble called `airports` (as well as others). Use these two dataframes to find the answer to 3 of the following, and print them out in a separate chunk (i.e. the chunk should print the tibble, thus showing the first 10 lines of each):

**Answer:**

First of all, let’s load the necessary libraries

```
library(tidyverse)
library(ggplot2)
library(nycflights13)
library(dplyr)
```

The number of flights (in the whole year) to each destination

```
flights %>%
  group_by(year, dest) %>%
  summarise(
    count = n()
  )
```

```
## # A tibble: 105 x 3
## # Groups:   year [1]
##   year dest count
##   <int> <chr> <int>
## 1  2013 ABQ    254
## 2  2013 ACK    265
## 3  2013 ALB    439
## 4  2013 ANC     8
## 5  2013 ATL  17215
## 6  2013 AUS  2439
## 7  2013 AVL    275
## 8  2013 BDL    443
## 9  2013 BGR    375
## 10 2013 BHM    297
## # ... with 95 more rows
```

The number and list of distinct airports in the US

- The number of distinct airports in the US:

```
airports %>%
  summarise(
```

```
nr_of_airports = n_distinct(name, na.rm = TRUE)
)
```

```
## # A tibble: 1 x 1
##   nr_of_airports
##           <int>
## 1           1440
```

- The list of distinct airports in the US along with how many times they are featured in the `airports` tibble:

```
airports %>%
  group_by(name) %>%
  summarise(
    nr_of_times_featured = n_distinct(name, na.rm = TRUE)
  )
```

```
## # A tibble: 1,440 x 2
##   name                                nr_of_times_featured
##   <chr>                                <int>
## 1 Aberdeen Regional Airport             1
## 2 Abilene Rgnl                         1
## 3 Abraham Lincoln Capital              1
## 4 Acadiana Rgnl                        1
## 5 Adak Airport                         1
## 6 Adams Fld                           1
## 7 Addison                             1
## 8 Adirondack Regional Airport           1
## 9 Akhiok Airport                       1
## 10 Akiak Airport                       1
## # ... with 1,430 more rows
```

The number and list of distinct airports that have at least one flight in the whole year from NYC

- The number of distinct airports that have at least one flight in the whole year from NYC

```
flights %>%
  filter(origin == 'JFK') %>%
  summarise(
    nr_of_times_featured = n_distinct(dest, na.rm = TRUE)
  )
```

```
## # A tibble: 1 x 1
##   nr_of_times_featured
##           <int>
## 1                70
```

- The list of distinct airports that have at least one flight in the whole year from NYC

```
flights %>%
  group_by(dest) %>%
  filter(origin == 'JFK') %>%
  summarise(
    nr_of_times_featured = n_distinct(dest, na.rm = TRUE)
  )
```

```
## # A tibble: 70 x 2
##   dest nr_of_times_featured
```

```
##      <chr>                <int>
## 1 ABQ                      1
## 2 ACK                      1
## 3 ATL                      1
## 4 AUS                      1
## 5 BHM                      1
## 6 BNA                      1
## 7 BOS                      1
## 8 BQN                      1
## 9 BTV                      1
## 10 BUF                     1
## # ... with 60 more rows
```

**Exercise 2:** Find all the rows with NA values in **the first two columns** for the following datasets:

- diamonds

```
diamonds %>%
  filter(is.na(carat) | is.na(cut))

## # A tibble: 0 x 10
## # ... with 10 variables: carat <dbl>, cut <ord>, color <ord>,
## #   clarity <ord>, depth <dbl>, table <dbl>, price <int>, x <dbl>,
## #   y <dbl>, z <dbl>
## # seems like there are no such rows
```

- flights

```
flights %>%
  filter(is.na(year) | is.na(month))

## # A tibble: 0 x 19
## # ... with 19 variables: year <int>, month <int>, day <int>,
## #   dep_time <int>, sched_dep_time <int>, dep_delay <dbl>, arr_time <int>,
## #   sched_arr_time <int>, arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>,
## #   distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
## # seems like there are no such rows
```

- mtcars

```
mtcars %>%
  filter(is.na(mpg) | is.na(cyl))

## [1] mpg cyl disp hp drat wt qsec vs am gear carb
## <0 rows> (or 0-length row.names)
## # seems like there are no such rows
```

The next exercise asks you to check *all* columns, but you don't want to do that with `filter()`. Why do you not want to do that with `filter`, especially for a dataset with hundreds of columns?

If there are lots of columns, we would need a very long expression with many repetitions of `is.na()` which is difficult to read.

**Exercise 3:** Look up `filter_all` and look at the examples at the end of the documentation. Use this (with some google-fu or discourse help) to find all the rows with NA values in *any* column for the following datasets:

Looking up `filter_all`:

```
?filter_all
```

- diamonds

```
filter_all(diamonds, any_vars(is.na(.)))
```

```
## # A tibble: 0 x 10
## #   ... with 10 variables: carat <dbl>, cut <ord>, color <ord>,
## #   clarity <ord>, depth <dbl>, table <dbl>, price <int>, x <dbl>,
## #   y <dbl>, z <dbl>
```

```
# seems like there are no such rows
```

- flights

```
filter_all(flights, any_vars(is.na(.)))
```

```
## # A tibble: 9,430 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1    1525           1530         -5    1934
## 2  2013     1     1    1528           1459         29    2002
## 3  2013     1     1    1740           1745         -5    2158
## 4  2013     1     1    1807           1738         29    2251
## 5  2013     1     1    1939           1840         59      29
## 6  2013     1     1    1952           1930         22   2358
## 7  2013     1     1    2016           1930         46     NA
## 8  2013     1     1      NA           1630         NA     NA
## 9  2013     1     1      NA           1935         NA     NA
## 10 2013     1     1      NA           1500         NA     NA
## # ... with 9,420 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

- mtcars

```
filter_all(mtcars, any_vars(is.na(.)))
```

```
## [1] mpg cyl disp hp drat wt  qsec vs  am  gear carb
## <0 rows> (or 0-length row.names)
```

```
# seems like there are no such rows
```

Thus, the output should be those rows that *do* contain NA values. Then look up `na.omit` (hat tip @kristof) and use that to achieve the same goal.

Looking up `na.omit`:

```
?na.omit
```

- diamonds

```
na.omit(diamonds, invert=TRUE)
```

```
## # A tibble: 53,940 x 10
##   carat cut      color clarity depth table price      x      y      z
##   <dbl> <ord>    <ord> <ord>   <dbl> <dbl> <int> <dbl> <dbl> <dbl>
## 1 0.23 Ideal    E     SI2    61.5   55   326  3.95  3.98  2.43
## 2 0.21 Premium E     SI1    59.8   61   326  3.89  3.84  2.31
## 3 0.23 Good    E     VS1    56.9   65   327  4.05  4.07  2.31
```

```
## 4 0.290 Premium I VS2 62.4 58 334 4.2 4.23 2.63
## 5 0.31 Good J SI2 63.3 58 335 4.34 4.35 2.75
## 6 0.24 Very Good J VVS2 62.8 57 336 3.94 3.96 2.48
## 7 0.24 Very Good I VVS1 62.3 57 336 3.95 3.98 2.47
## 8 0.26 Very Good H SI1 61.9 55 337 4.07 4.11 2.53
## 9 0.22 Fair E VS2 65.1 61 337 3.87 3.78 2.49
## 10 0.23 Very Good H VS1 59.4 61 338 4 4.05 2.39
## # ... with 53,930 more rows
```

- flights

```
na.omit(flights, invert = TRUE)
```

```
## # A tibble: 327,346 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1 2013     1     1     517             515           2     830
## 2 2013     1     1     533             529           4     850
## 3 2013     1     1     542             540           2     923
## 4 2013     1     1     544             545          -1    1004
## 5 2013     1     1     554             600          -6     812
## 6 2013     1     1     554             558          -4     740
## 7 2013     1     1     555             600          -5     913
## 8 2013     1     1     557             600          -3     709
## 9 2013     1     1     557             600          -3     838
## 10 2013     1     1     558             600          -2     753
## # ... with 327,336 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

- mtcars

```
na.omit(mtcars, invert = TRUE)
```

```
##           mpg cyl  disp  hp drat   wt  qsec vs am gear carb
## Mazda RX4      21.0   6 160.0 110 3.90 2.620 16.46 0 1 4 4
## Mazda RX4 Wag  21.0   6 160.0 110 3.90 2.875 17.02 0 1 4 4
## Datsun 710      22.8   4 108.0  93 3.85 2.320 18.61 1 1 4 1
## Hornet 4 Drive  21.4   6 258.0 110 3.08 3.215 19.44 1 0 3 1
## Hornet Sportabout 18.7   8 360.0 175 3.15 3.440 17.02 0 0 3 2
## Valiant        18.1   6 225.0 105 2.76 3.460 20.22 1 0 3 1
## Duster 360     14.3   8 360.0 245 3.21 3.570 15.84 0 0 3 4
## Merc 240D      24.4   4 146.7  62 3.69 3.190 20.00 1 0 4 2
## Merc 230       22.8   4 140.8  95 3.92 3.150 22.90 1 0 4 2
## Merc 280       19.2   6 167.6 123 3.92 3.440 18.30 1 0 4 4
## Merc 280C      17.8   6 167.6 123 3.92 3.440 18.90 1 0 4 4
## Merc 450SE     16.4   8 275.8 180 3.07 4.070 17.40 0 0 3 3
## Merc 450SL     17.3   8 275.8 180 3.07 3.730 17.60 0 0 3 3
## Merc 450SLC    15.2   8 275.8 180 3.07 3.780 18.00 0 0 3 3
## Cadillac Fleetwood 10.4   8 472.0 205 2.93 5.250 17.98 0 0 3 4
## Lincoln Continental 10.4   8 460.0 215 3.00 5.424 17.82 0 0 3 4
## Chrysler Imperial 14.7   8 440.0 230 3.23 5.345 17.42 0 0 3 4
## Fiat 128       32.4   4  78.7  66 4.08 2.200 19.47 1 1 4 1
## Honda Civic    30.4   4  75.7  52 4.93 1.615 18.52 1 1 4 2
## Toyota Corolla 33.9   4  71.1  65 4.22 1.835 19.90 1 1 4 1
```

## Toyota Corona	21.5	4	120.1	97	3.70	2.465	20.01	1	0	3	1
## Dodge Challenger	15.5	8	318.0	150	2.76	3.520	16.87	0	0	3	2
## AMC Javelin	15.2	8	304.0	150	3.15	3.435	17.30	0	0	3	2
## Camaro Z28	13.3	8	350.0	245	3.73	3.840	15.41	0	0	3	4
## Pontiac Firebird	19.2	8	400.0	175	3.08	3.845	17.05	0	0	3	2
## Fiat X1-9	27.3	4	79.0	66	4.08	1.935	18.90	1	1	4	1
## Porsche 914-2	26.0	4	120.3	91	4.43	2.140	16.70	0	1	5	2
## Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2
## Ford Pantera L	15.8	8	351.0	264	4.22	3.170	14.50	0	1	5	4
## Ferrari Dino	19.7	6	145.0	175	3.62	2.770	15.50	0	1	5	6
## Maserati Bora	15.0	8	301.0	335	3.54	3.570	14.60	0	1	5	8
## Volvo 142E	21.4	4	121.0	109	4.11	2.780	18.60	1	1	4	2

The argument `invert` does not seem to work (if set to `TRUE`, only the rows with NA should remain). The reason for this according to comments of people on several forums is that in this version of R `na.omit()` does not have the argument `invert` anymore. As simply negating the function with `!` does not work either, I could not figure out how to keep only the rows with NA values using `na.omit()`.

**Exercise 5:** Come up with an exercise to help you – and others – learn `summarise` and `group_by` better. The more confused you are, the more you should simply try to come up with, or even copy, an example from somewhere and highlight what confuses you. Is it the order or arguments? Their role? If you are less confused, try to find a (non-obvious) use. Mention any resources used.

**Answer:** How much time did those planes (identified by tailnum) gain on average in air in a month which took off at least 30 minutes late but arrived earlier than scheduled time or on time?

```
flights %>%
  group_by(year, month, tailnum) %>%
  filter(dep_delay >= 30.0, arr_delay <= 0) %>%
  summarise(gain_time_mean = mean(dep_delay, na.rm = TRUE) - mean(arr_delay, na.rm = TRUE))
```

```
## # A tibble: 519 x 4
## # Groups:   year, month [12]
##   year month tailnum gain_time_mean
##   <int> <int> <chr>          <dbl>
## 1  2013     1 N369NW             43
## 2  2013     1 N3730B             50
## 3  2013     1 N377DA             45
## 4  2013     1 N3BJAA             34
## 5  2013     1 N436UA             31
## 6  2013     1 N437UA             38
## 7  2013     1 N438UA             36
## 8  2013     1 N4WAAA             45
## 9  2013     1 N524JB             37
## 10 2013     1 N539UA             56
## # ... with 509 more rows
```

**Exercise 6:** Work through sections 11.1 and 11.2 (skip exercises).

**Answer:**

### 11.1 Prerequisites

```
library(tidyverse)
```

### 11.2 Getting started

Reading a csv

```
testdata3 <- read_csv("D:/Egyetem/CEU/Coding_1/R-Coding/lecture6/test-data3.csv")
```

```
(testdata3)
```

```
## # A tibble: 3,009 x 5
##   amount also_amount who_knows name has_passed
##   <dbl>      <dbl>    <dbl> <chr> <lgl>
## 1     1         2      3 a    TRUE
## 2     4         5      6 b    TRUE
## 3     7         8      9 c    FALSE
## 4     1         2      3 a    TRUE
## 5     4         5      6 b    TRUE
## 6     7         8      9 c    FALSE
## 7     7         8      9 c    FALSE
## 8     7         8      9 c    FALSE
## 9     7         8      9 c    FALSE
## 10    7         8      9 c    FALSE
## # ... with 2,999 more rows
```

Supplying inline csv

```
read_csv("a,b,c
1,2,3
4,5,6")
```

```
## # A tibble: 2 x 3
##       a     b     c
##   <dbl> <dbl> <dbl>
## 1     1     2     3
## 2     4     5     6
```

Skipping the first n lines

```
read_csv("First line of metadata goes here
Second line of metadata goes here
x,y,z
1,2,3", skip = 2)
```

```
## # A tibble: 1 x 3
##       x     y     z
##   <dbl> <dbl> <dbl>
## 1     1     2     3
```

Comments to skip

```
read_csv("# A comment I want to skip
x,y,z
1,2,3", comment = "#")
```

```
## # A tibble: 1 x 3
##       x     y     z
##   <dbl> <dbl> <dbl>
## 1     1     2     3
```

When there are no column names

```
read_csv("1,2,3\n4,5,6", col_names = FALSE)
```

```
## # A tibble: 2 x 3
```

```
##      X1      X2      X3
## <dbl> <dbl> <dbl>
## 1      1      2      3
## 2      4      5      6
```