

Assignment 7

Veronika Palotai

11/2/2019

Exercise 1

All done although setting up an existing directory as a project was a bit different from creating a completely new one. The necessary steps were 'File' -> 'New Project' -> 'Existing Directory'.

Exercise 2

All done.

Exercise 3

Import Libraries

```
library(ggplot2)
library(tidyverse)
library(scales)
```

- Import from csv

Setting the Path

```
dir <- "D:/Egyetem/CEU/Coding_1/R-Coding/"
```

Location Folders

```
data_in <- paste0(dir, "data_repo/hotels-europe/clean/")
```

Loading Datasets

```
hotels_prices_csv <- read_csv(paste0(data_in, "hotels-europe_price.csv"))
hotels_features_csv <- read_csv(paste0(data_in, "hotels-europe_features.csv"))
```

- Import from .dta

Import 'haven' library and look up read_dta

```
library(haven)
?read_dta
```

Reading .dta files

```
hotels_prices_dta <- read_dta(paste0(data_in, "hotels-europe_price.dta"), encoding = NULL)
hotels_features_dta <- read_dta(paste0(data_in, "hotels-europe_features.dta"), encoding = NULL)
```

- Comparing contents

First, let's look at the contents of the tables.

```
View(hotels_features_dta)
View(hotels_features_csv)
View(hotels_prices_dta)
View(hotels_prices_csv)
```

We can see that the order of the fields and the records are different. Let's check the length of each table to see whether they contain equal number of observations.

```
hotels_features_csv %>%
  summarise(
    nr_of_records = n()
  )
```

```
## # A tibble: 1 x 1
##   nr_of_records
##         <int>
## 1         22902
```

```
hotels_features_dta %>%
  summarise(
    nr_of_records = n()
  )
```

```
## # A tibble: 1 x 1
##   nr_of_records
##         <int>
## 1         22902
```

```
hotels_prices_csv %>%
  summarise(
    nr_of_records = n()
  )
```

```
## # A tibble: 1 x 1
##   nr_of_records
##         <int>
## 1         148021
```

```
hotels_prices_dta %>%
  summarise(
    nr_of_records = n()
  )
```

```
## # A tibble: 1 x 1
##   nr_of_records
##         <int>
## 1         148021
```

As we can see, they contain the same number of records. Let's go into the details a bit and check the number of records for each city in the features table and each year in the prices table.

```
hotels_features_csv %>%
  group_by(city) %>%
  summarise(
    nr_of_records = n()
  )
```

```
## # A tibble: 47 x 2
##   city      nr_of_records
##   <chr>          <int>
## 1 Amsterdam         428
## 2 Athens            369
## 3 Barcelona         835
## 4 Belgrade          111
## 5 Berlin            579
## 6 Birmingham         185
## 7 Bratislava        104
## 8 Brussels          333
## 9 Bucharest         132
## 10 Budapest         340
## # ... with 37 more rows
```

```
hotels_features_dta %>%
  group_by(city) %>%
  summarise(
    nr_of_records = n()
  )
```

```
## # A tibble: 46 x 2
##   city      nr_of_records
##   <chr>          <int>
## 1 Amsterdam         429
## 2 Athens            369
## 3 Barcelona         835
## 4 Belgrade          111
## 5 Berlin            579
## 6 Birmingham         185
## 7 Bratislava        104
## 8 Brussels          333
## 9 Bucharest         132
## 10 Budapest         340
## # ... with 36 more rows
```

```
hotels_prices_csv %>%
  group_by(year) %>%
  summarise(
    nr_of_records = n()
  )
```

```
## # A tibble: 2 x 2
##   year nr_of_records
##   <dbl>          <int>
## 1  2017         62235
## 2  2018         85786
```

```
hotels_prices_dta %>%
  group_by(year) %>%
  summarise(
    nr_of_records = n()
  )
```

```
## # A tibble: 2 x 2
##   year nr_of_records
##   <dbl>         <int>
## 1  2017         62235
## 2  2018         85786
```

Most of the time they are the same in both tables, however, we can see a difference in case of ‘Amsterdam’ which might be due to a reading error. Since the overall number of records are the same, there has to be at least one other city with a difference.

Exercise 4

Picking the first 200 lines of the prices table

```
(hotels_prices_csv_small <- hotels_prices_csv %>%
  slice(1:200))
```

```
## # A tibble: 200 x 10
##   hotel_id price offer offer_cat year month weekend holiday nnights
##   <dbl> <dbl> <dbl> <chr>    <dbl> <dbl> <dbl> <dbl> <dbl>
## 1      1    172     0 0% no of~  2017    11     1     0     1
## 2      1    122     1 15-50% o~  2018     1     1     0     1
## 3      1    122     1 15-50% o~  2017    12     0     1     1
## 4      1    552     1 1-15% of~  2017    12     0     1     4
## 5      1    122     1 15-50% o~  2018     2     1     0     1
## 6      1    114     1 15-50% o~  2017    11     0     0     1
## 7      2    119     0 0% no of~  2017    11     0     0     1
## 8      2    119     0 0% no of~  2017    12     0     1     1
## 9      2    547     0 0% no of~  2017    12     0     1     4
## 10     3    118     1 15-50% o~  2017    12     0     1     1
## # ... with 190 more rows, and 1 more variable: scarce_room <dbl>
```

Writing to file

```
write.csv(hotels_prices_csv_small, file = "hotels-europe-small.csv", row.names = FALSE)
```

- Lines changed: I deleted the field names
- It lead to the problem that the contents of the first row were interpreted as column names
- Fixing:

```
hotels_small <- read.csv("hotels-europe-small.csv", header = FALSE)
# names columns as V1, V2, V3, ...
```

Exercise 5

12.6 Case Study

Load the dataset:

```
who <- tidyr::who
```

Gather together all the columns from `new_sp_m014` to `newrel_f65` because they are values not variables:

```
who1 <- who %>%  
  gather(new_sp_m014:newrel_f65, key = "key", value = "cases", na.rm = TRUE)
```

Some hint on the structure of the values in the new key column by counting them:

```
who1 %>%  
  count(key)
```

```
## # A tibble: 56 x 2  
##   key          n  
##   <chr>      <int>  
## 1 new_ep_f014  1032  
## 2 new_ep_f1524 1021  
## 3 new_ep_f2534 1021  
## 4 new_ep_f3544 1021  
## 5 new_ep_f4554 1017  
## 6 new_ep_f5564 1017  
## 7 new_ep_f65   1014  
## 8 new_ep_m014  1038  
## 9 new_ep_m1524 1026  
## 10 new_ep_m2534 1020  
## # ... with 46 more rows
```

Minor fix to the format of the column names:

```
who2 <- who1 %>%  
  mutate(key = stringr::str_replace(key, "newrel", "new_rel"))
```

Separate the values in each code with two passes of `separate()`. The first pass will split the codes at each underscore.

```
who3 <- who2 %>%  
  separate(key, c("new", "type", "sexage"), sep = "_")
```

Drop the 'new' column because it's constant in this dataset. Let's also drop 'iso2' and 'iso3' since they're redundant.

```
who3 %>%  
  count(new)
```

```
## # A tibble: 1 x 2  
##   new          n  
##   <chr> <int>  
## 1 new    76046
```

```
who4 <- who3 %>%
  select(-new, -iso2, -iso3)
```

Separate 'sexage' into 'sex' and 'age' by splitting after the first character:

```
who5 <- who4 %>%
  separate(sexage, c("sex", "age"), sep = 1)
```

The 'who' dataset is now tidy, however, we could have done all these steps at the same time, like this:

```
who %>%
  gather(key, value, new_sp_m014:newrel_f65, na.rm = TRUE) %>%
  mutate(key = stringr::str_replace(key, "newrel", "new_rel")) %>%
  separate(key, c("new", "var", "sexage")) %>%
  select(-new, -iso2, -iso3) %>%
  separate(sexage, c("sex", "age"), sep = 1)
```

```
## # A tibble: 76,046 x 6
##   country      year var   sex   age  value
##   <chr>      <int> <chr> <chr> <chr> <int>
## 1 Afghanistan 1997 sp    m    014     0
## 2 Afghanistan 1998 sp    m    014    30
## 3 Afghanistan 1999 sp    m    014     8
## 4 Afghanistan 2000 sp    m    014    52
## 5 Afghanistan 2001 sp    m    014   129
## 6 Afghanistan 2002 sp    m    014    90
## 7 Afghanistan 2003 sp    m    014   127
## 8 Afghanistan 2004 sp    m    014   139
## 9 Afghanistan 2005 sp    m    014   151
## 10 Afghanistan 2006 sp    m    014   193
## # ... with 76,036 more rows
```

Exercise 6

The dataset:

```
(df <- tibble(name = c("A123", "B456"), age = c(30, 60), answer1 = c(0, 1), answer2 = c(1,1), answer3 =
```

```
## # A tibble: 2 x 6
##   name      age answer1 answer2 answer3 answer4
##   <chr> <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1 A123    30         0         1         1         0
## 2 B456    60         1         1         0         0
```

1. Cleaning by hand - description only

What needs to be done is to gather the columns 'answers1' to 'answers4' and create an 'answer_nr' (key) and an 'answer_value' (value) column for the answers and their corresponding values pertaining to a specific age group.

2. Cleaning by coding it with `gather()`

```
tidy_df <- df %>%
  gather(answer1:answer4, key="answer_nr", value="answer_value")
tidy_df
```

```
## # A tibble: 8 x 4
##   name    age answer_nr answer_value
##   <chr> <dbl> <chr>         <dbl>
## 1 A123    30 answer1           0
## 2 B456    60 answer1           1
## 3 A123    30 answer2           1
## 4 B456    60 answer2           1
## 5 A123    30 answer3           1
## 6 B456    60 answer3           0
## 7 A123    30 answer4           0
## 8 B456    60 answer4           0
```